

Safety and Reliability

1. Input Guardrails

Validates user queries for harmful or invalid content before processing.

2. PII Redaction

Automatically removes personal information from data retrieved from Gmail, Notion, and Jira.

3. Citation Validation

Every answer claim is linked to actual retrieved data to prevent hallucination.

4. Output Validation

Final answers are checked for accuracy, format, and safety before delivery.

5. RAG Evaluation

Regular testing with normal, tricky, and missing data queries to measure system performance.

Limitations

1. Real Enterprise Data Access

We can only use test accounts and sample data, not actual company data with real employees.

2. API Authentication Scope

Limited to test accounts for Gmail, Notion, and Jira due to app verification requirements.

3. API Rate Limits

Free/developer API tiers impose request limits, unlike paid enterprise plans.

4. Real PII for Testing


We can only test PII redaction with synthetic data, not genuine sensitive information.

5. Tool Selection at Scale

Tool selection accuracy degrades beyond about a dozen tools (drops to 13.6% baseline, recovers to 43% with retrieval filtering). Not a constraint for our 3-server MVP, but a real limit when scaling.

System Workflows

MVP LangGraph Workflow



Memory checkpointer and LangSmith tracing run across every node

Permission Layer (v2)

